

XDP (eXpress Data Path) as a building block for other FOSS projects

Jesper Dangaard Brouer (Red Hat)

LLC (Lund Linux Con)
Lund, Sweden, May 2019

What will you learn?

Contents of this talk

Framing XDP

- Why it is needed? (and how is it different and better)
- How it is a **building block**? (evolving and use-cases I hope to see)

Framing AF_XDP

- Explain: **What is AF_XDP?**
- How XDP (redirect) is building block for AF_XDP
- Explaining where AF_XDP performance comes from.

What is XDP?

XDP (eXpress Data Path) is a Linux **in-kernel** fast-path

- **New programmable layer in-front** of traditional network stack
- Already accepted part of upstream kernels (and RHEL8)
- Operate at the same level and speeds as DPDK
- For L2-L3 use-cases: seeing x10 performance improvements!
- Can accelerate **in-kernel** L2-L3 use-cases (e.g. forwarding)
- Force NIC driver to register their **memory model**
 - when enabling XDP **driver MUST switch to supported memory model**

What is AF_XDP?

XDP is **in-kernel** processing... how to get frames into userspace?

What is **AF_XDP**? (the Address Family XDP socket)

- Hybrid **kernel-bypass** facility, move selective frames out of kernel
- One user of the **XDP_REDIRECT** feature
 - XDP/eBPF prog filters packets using REDIRECT into AF_XDP socket
- Delivers raw **L2 frames into userspace**

AF_XDP: Copy vs. Zero-Copy mode

- Drivers with XDP_REDIRECT have 1-copy into userspace memory
- ZC: Drivers need **zero-copy memory model** (and XDP_SETUP_XSK_UMEM)
 - Allow: use userspace memory directly in NIC DMA ring

Why is XDP needed?

This is about: **Kernel networking stack staying relevant**

- For emerging use-cases and areas

Linux networking stack optimized for layers L4-L7

- Missing something to address L2-L3 use-cases

XDP operate at layers L2-L3

If you forgot OSI model:

- L2=Ethernet
- L3=IPv4/IPv6
- L4=TCP/UDP
- L7=Applications

Existing solutions: Not first mover

XDP is not first mover in this area

- But we believe **XDP is different and better**

Existing **kernel bypass** solutions:

- netmap (FreeBSD), DPDK (Intel/LF), PF_ring (ntop)
- maglev (Google), Onload (SolarFlare), Snabb

Commercial solutions **similar to XDP**:

- ndiv by HAproxy, product **ALOHA**

What makes XDP different and better?

Not bypass, but in-kernel fast-path

The killer feature of XDP is integration with Linux kernel,

- Leverages existing kernel drivers and infrastructure
- Programmable flexibility via eBPF sandboxing (kernel infra)
- Flexible sharing of NIC resources between Linux and XDP
- Cooperation with netstack via eBPF-helpers and fallback-handling
- No need to reinject packets (unlike bypass solutions)

AF_XDP for flexible kernel bypass

- Cooperate with use-cases needing fast raw frame access in userspace
- While leveraging existing kernel NIC drivers

XDP is a building block

Fundamental to understand that XDP is a building block

XDP is a building block

It is fundamental to understand

XDP is a **component**; a core facility provided by the kernel

- Put it together with other components to solve a task

eBPF (incl XDP) is **not a product in itself**

- Existing (and new) Open Source projects will use these eBPF components

Full potential comes when

- Combining XDP-eBPF with other eBPF-hooks and facilities
- To construct a **"networking pipeline"** via kernel components
- The **Cilium** project is a good example (container L4-L7 policy)

XDP use-cases

Areas and use-cases where XDP is **already being used**

Touch upon **new potential** and opportunities

- e.g. for Virtual Machines (VM) and Containers

Use-case: Anti-DDoS

The most obvious use case for XDP is **anti-DDoS**

Companies already deployed XDP in production for anti-DDoS

- **Facebook**, every packet goes through XDP for **1.5 years**
- **CloudFlare** switched to **XDP** (changed NIC vendor due to XDP support!)

New potential: Protecting Containers and VMs

- **Containers:** Protect Kubernetes/OpenShift cluster with XDP
- **VM:** Host-OS protect Guest-OS'es via XDP
 - Work-in-progress: allow vhost/virtio_net; upload XDP to Host-OS

Use-case: L4 Load-balancer

Facebook was using the kernel Load-balancer IPVS

- Switched to using XDP instead: Reported **x10 performance improvement**
- Open Sourced their **XDP load-balancer** called **katran**

New potential: Host OS load-balancing to VMs and Containers

- **VM:** Phy-NIC can XDP_REDIRECT into Guest-NIC
 - driver tuntap queues XDP-raw frames to virtio_net; **skip SKB in Host-OS**
- **Container:** Phy-NIC can XDP_REDIRECT into **veth** (kernel v4.20)
 - driver veth allocs+builds SKB outside driver-code; speedup **skip some code**
 - veth can **RE-redirect**, allow **building interesting proxy-solutions**

Evolving XDP via leveraging existing solutions

XDP can (easily) be misused in the same way as kernel bypass solutions

Being smart about how XDP is integrated into existing Open Source solutions

- Leverage existing eco-systems e.g. for control plane setup

Evolving XDP via BPF-helpers

We should encourage adding helpers instead of duplicating data in BPF maps

Think of XDP as a **software offload layer for the kernel netstack**

- Simply setup and use the Linux netstack, but accelerate parts of it with XDP

IP routing good example: **Access routing table from XDP via BPF helpers** (v4.18)

- Let Linux handle routing (daemons) and neighbour lookups
- Talk at LPC-2018 (David Ahern): **Leveraging Kernel Tables with XDP**

Obvious **next target**: **Bridge lookup helper**

- Like IP routing: transparent XDP acceleration of bridge forwarding
 - Fallback for ARP lookups, flooding etc.
- Huge potential **performance boost for Linux bridge** use cases!

Transfer info between XDP and netstack

Ways to transfer info between XDP and netstack

- XDP can modify packet headers before netstack
 - Pop/push headers influence RX-handler in netstack
 - CloudFlare modifies MAC-src on sampled dropped packets
- XDP have 32 bytes metadata in front of payload
 - TC eBPF (cls_bpf) can read this, and update SKB fields
 - E.g. save XDP lookup and use in TC eBPF hook
 - AF_XDP raw frames have this metadata avail in front of payload

XDP integration with OVS

XDP/eBPF can integrate/offload Open vSwitch (OVS) in many ways

- VMware (William Tu) presented different options at LPC 2018:
 - Bringing the Power of eBPF to Open vSwitch
- TC eBPF, (re)implemented OVS in eBPF (performance limited)
- Offloading subset to XDP (issue: missing some BPF helpers)
- AF_XDP, huge performance gain

AF_XDP – moving frames into userspace

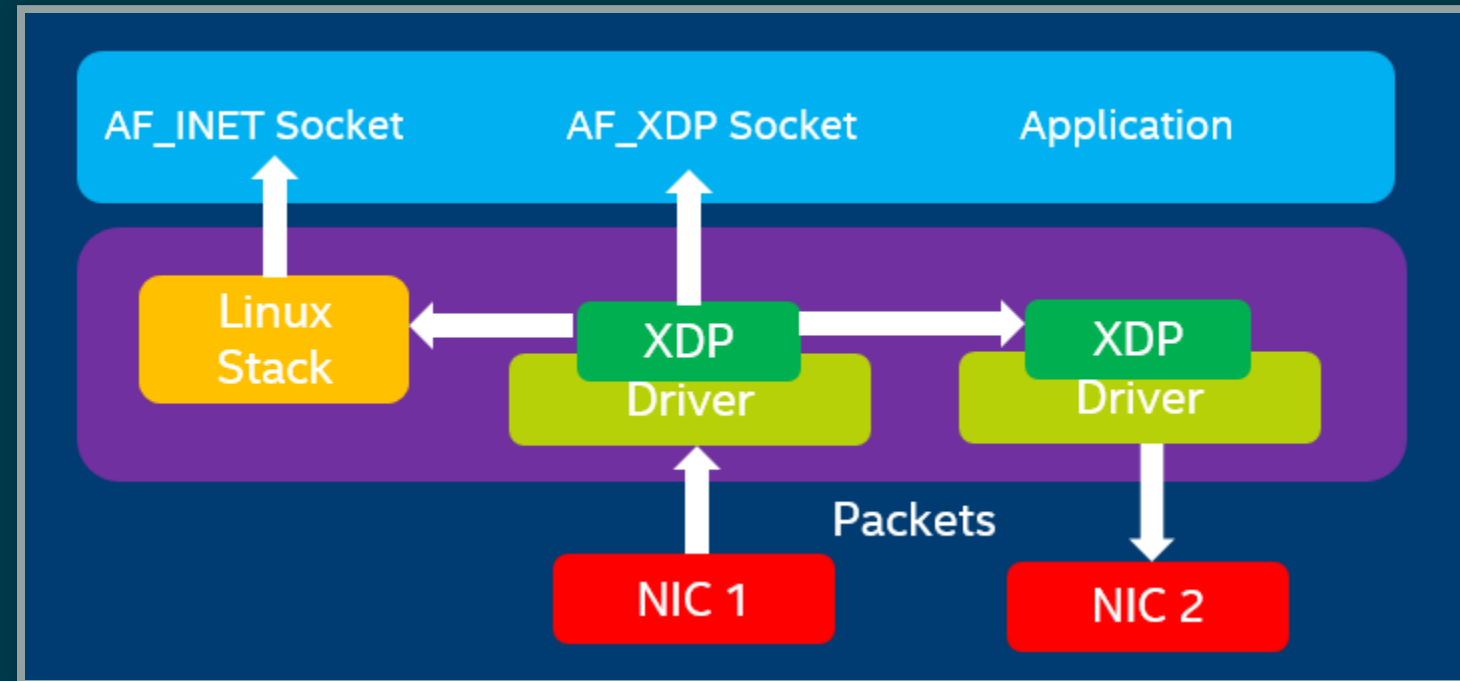
XDP is in-kernel processing

AF_XDP is hybrid kernel-bypass facility

- XDP filter and redirect raw frames into userspace

Userspace AF_XDP socket pre-register user memory

AF_XDP Basics



XDP program filters packets

- for AF_XDP socket redirect (fast-path packets)
- or normal processing by Linux network stack

Integration with AF_XDP

How can **kernel-bypass** solutions use AF_XDP as a **building block**?

AF_XDP integration with DPDK

AF_XDP poll-mode driver for DPDK

- DPDK v19.05-rc1 integrates AF_XDP Poll-Mode-Driver
- ~1% overhead

Advantages:

- Don't monopolize entire NIC
- Split traffic to kernel with XDP filter program
- HW independent application binary
- Isolation and robustness
- Cloud-native support
- Fewer setup restrictions

AF_XDP integration with VPP

VPP (FD.io) **could** integrate via AF_XDP DPDK PMD

- But VPP uses only user-mode driver of DPDK
- VPP has a lot of native functionality

A native AF_XDP driver would be more efficient

- Less code and easier setup without DPDK

AF_XDP integration with Snabb Switch

Snabb Switch

- Implement an **AF_XDP driver?**
- Allow leveraging kernel drivers that implement XDP
 - Kernel community takes care of maintaining driver code
- Any **performance loss/gap** to native Snabb driver ?
 - E.g. NAPI “only” bulk up-to 64 packets
 - E.g. NAPI is not doing busy-polling 100%, more latency variance

Explaining AF_XDP performance

Where does AF_XDP performance come from?

Lock-free channel directly from driver RX-queue into AF_XDP socket

- Single-Producer/Single-Consumer (SPSC) descriptor ring queues
- Single-Producer (SP) via bind to specific RX-queue id
 - NAPI-softirq assures only 1-CPU process 1-RX-queue id (per sched)
- Single-Consumer (SC) via 1-Application
- Bounded buffer pool (UMEM) allocated by userspace (register with kernel)
 - Descriptor(s) in ring(s) point into UMEM
 - No memory allocation, but return frames to UMEM in timely manner
- Transport signature Van Jacobson talked about
 - Replaced by XDP/eBPF program choosing to XDP_REDIRECT

Details: Actually **four** SPSC ring queues

AF_XDP **socket**: Has **two rings**: **RX** and **TX**

- Descriptor(s) in ring points into UMEM

UMEM consists of a number of equally sized chunks

- Has **two rings**: **FILL** ring and **COMPLETION** ring
- FILL ring: application gives kernel area to RX fill
- COMPLETION ring: kernel tells app TX is done for area (can be reused)

Gotcha by RX-queue id binding

AF_XDP bound to **single RX-queue id** (for SPSC performance reasons)

- NIC by default spreads flows with RSS-hashing over RX-queues
 - Traffic likely not hitting queue you expect
- You **MUST** configure NIC **HW filters** to **steer to RX-queue id**
 - Out of scope for XDP setup
 - Use ethtool or TC HW offloading for filter setup
- **Alternative** work-around
 - Create as many AF_XDP sockets as RXQs
 - Have userspace poll()/select on all sockets

Summary

- XDP = Linux kernel fast path
- AF_XDP = packets to userspace from XDP
- Similar speeds as DPDK
- A building block for a solution. Not a ready solution in itself.
- Many upcoming use cases,
 - e.g., OVS, XDP-offload netstack, DPDK PMD
- Come join the fun!
 - <https://github.com/xdp-project/>
- XDP-tutorial at:
 - <https://github.com/xdp-project/xdp-tutorial/>