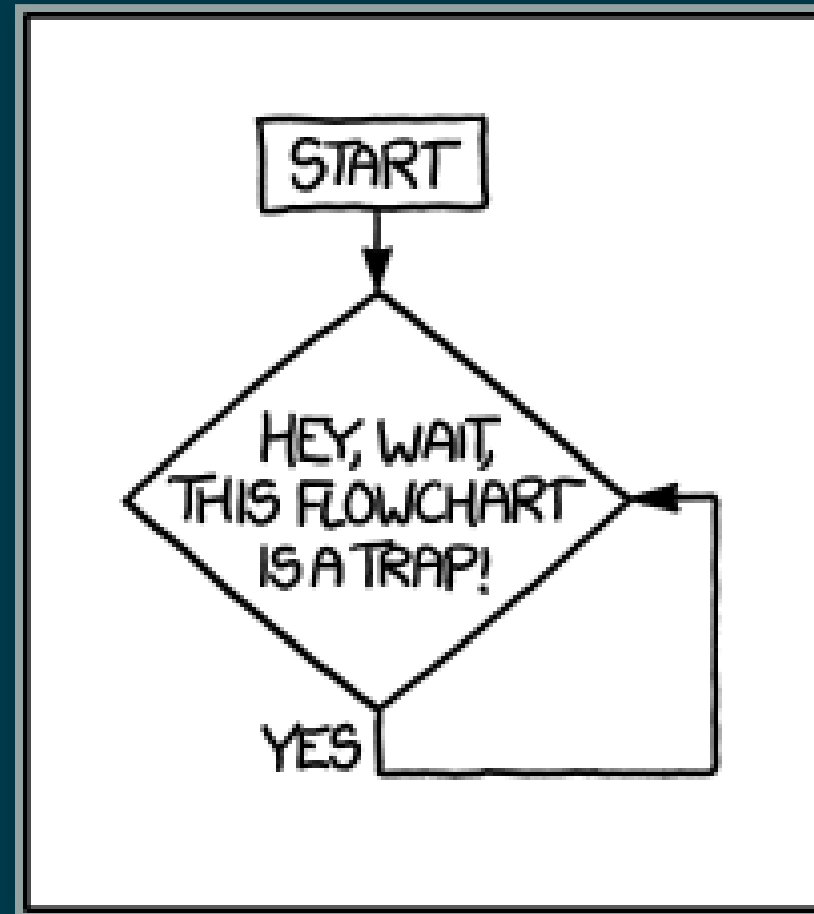


# XDP + BPF\_PROG\_TEST\_RUN = A programmable traffic generator!

Toke Høiland-Jørgensen  
Principal Kernel Engineer, Red Hat

Lund Linux Con  
April 2022

# Reminder: How does XDP work?



<https://xkcd.com/1195>

Hopefully you paid attention during the two previous talks :)

# What is BPF\_PROG\_TEST\_RUN?

```
struct { /* anonymous struct used by BPF_PROG_TEST_RUN command */
    __u32      prog_fd;
    __u32      retval;
    __u32      data_size_in; /* input: len of data_in */
    __u32      data_size_out; /* input/output: len of data_out
    *      returns ENOSPC if data_out
    *      is too small.
    */

    __aligned_u64 data_in;
    __aligned_u64 data_out;
    __u32      repeat;
    __u32      duration;
    __u32      ctx_size_in; /* input: len of ctx_in */
    __u32      ctx_size_out; /* input/output: len of ctx_out
    *      returns ENOSPC if ctx_out
    *      is too small.
    */

    __aligned_u64 ctx_in;
    __aligned_u64 ctx_out;
    __u32      flags;
    __u32      cpu;
} test;
```

# Introducing BPF\_F\_TEST\_XDP\_LIVE\_FRAMES

# Introducing BPF\_F\_TEST\_XDP\_LIVE\_FRAMES

Instead of returning the XDP program result to userspace, act on it!

- If XDP\_PASS, **inject** the packet into the stack (at `netif_receive_skb()`)
- If XDP\_REDIRECT or XDP\_TX, send packet **out** an interface (using `ndo_xdp_xmit()`)
- Optimise for batching using the `repeat` parameter (pages are recycled)
  - **Note!** Page recycling -> not the same data every time!

Comes with documentation! [https://docs.kernel.org/bpf/bpf\\_prog\\_run.html](https://docs.kernel.org/bpf/bpf_prog_run.html)

# What can we do with this?

- Packet injection into the kernel (using `XDP_PASS`)
- Testing `XDP_REDIRECT` infrastructure (like maps)
- Sending packets - XDP-powered traffic generator!

# The XDP traffic generator

Live demo time!

# The XDP traffic generator

In case the live demo didn't work:

```
# ./xdp-trafficgen udp ens3f1 # single core
Transmitting on ens3f1 (ifindex 6)
[..]
XDP_REDIRECT      11150720 pkts (   8919659 pps)      696920 KiB (   4567 Mbits/s)

# ./xdp-trafficgen udp ens3f1 -t 6 # 6 cores
Transmitting on ens3f1 (ifindex 6)
[..]
XDP_REDIRECT      65123603 pkts (  52095122 pps)     4070225 KiB ( 26673 Mbits/s)

# ./xdp-trafficgen udp ens3f1 -t 6 -d 100 # spraying over 100 dst-ports
Transmitting on ens3f1 (ifindex 6)
[..]
XDP_REDIRECT      8226576 pkts (  32896120 pps)     514161 KiB ( 16843 Mbits/s)
```



# How does it work?

- Prepare packet buffer in userspace
- Pass it to BPF\_PROG\_TEST\_RUN
- Return XDP\_REDIRECT with the right interface
  - For dynamic ports, update the port+checksum first

The simplest case is literally just:

```
SEC ("xdp")
int xdp_redirect_notouch(struct xdp_md *ctx)
{
    return xdp_stats_record_action(ctx, bpf_redirect(config.ifindex_out, 0));
}
```

# Can we do TCP as well?

Yes!

- Set up connection from userspace
- Install XDP program on interface to **intercept** replies
  - Process ACKs in XDP program, update state and drop packets
- Use BPF\_PROG\_TEST\_RUN to send out TCP packets
  - Return XDP\_DROP when running up against the CWND

```
Connected to fe80::ee0d:9aff:fedb:11cd port 1234 from fe80::ee0d:9aff:fed8:f5d3 port 39500
[...]
Period of 1.000081s ending at 1652215664.276825
XDP_DROP      3249504 pkts (    23878 pps)      4760015 KiB (   287 Mbits/s)
XDP_PASS      0 pkts (         0 pps)           0 KiB (         0 Mbits/s)
XDP_TX        516392331 pkts (  6217106 pps)    756434078 KiB ( 74605 Mbits/s) <--- retransmissions!
XDP_REDIRECT  545096864 pkts (  6534151 pps)    798481734 KiB ( 78410 Mbits/s)
```

Raw throughput (single core): **78.4 Gbps!** Goodput: **~4 Gbps!**

# End: Questions?

Why not try out XDP live packet mode yourself? What else can we use it for?

XDP-tools: [github.com/xdp-project/xdp-tools](https://github.com/xdp-project/xdp-tools)