

# XDP hints via BPF Type Format (BTF) system

Jesper Dangaard Brouer  
Senior Principal Kernel Engineer

Lund Linux Con  
April 2022

# Reminder: What is BPF ?

From: <https://ebpf.io/what-is-ebpf>

*BPF is a revolutionary technology that can run sandboxed programs in the Linux kernel without changing kernel source code or loading a kernel module*

BPF is a **technology name**: no longer an acronym

**Rate of innovation** at the operating system level: **Traditionally slow**

- BPF enables things at OS-level that were not possible before
- BPF will **radically increase** rate of innovation

# Traditional Kernel development process


Application Developer:

I want this new feature to observe my app



Hey kernel developer! Please add this new feature to the Linux kernel

OK! Just give me a year to convince the entire community that this is good for everyone.




1 year later...

I'm done. The upstream kernel now supports this.



But I need this in my Linux distro



5 year later...

Good news. Our Linux distribution now ships a kernel with your required feature

OK but my requirements have changed since...



# BPF development process

## Application Developer:

I want this new feature to observe my app



## eBPF Developer:

OK! The kernel can't do this so let me quickly solve this with eBPF.



## A couple of days later...

Here is a release of our eBPF project that has this feature now. BTW, you don't have to reboot your machine.





# Reminder: What is XDP?

XDP (eXpress Data Path) is a Linux **in-kernel** fast-path

- **Programmable layer in-front** of traditional network stack
  - Read, modify, drop, redirect or pass
  - For L2-L3 use-cases: seeing **x10 performance** improvements!
- **Avoiding memory allocations**
  - No SKB allocations and no-init (SKB zeroes 4 cache-lines per pkt)
- Adaptive **bulk** processing of frames
- Very **early access** to frame (in driver code **after DMA sync**)
- Ability to **skip (large parts) of kernel code**
  - Evolve XDP via **BPF-helpers**
- **Pitfall: Lost traditional HW offloads** (e.g. RX-hash, checksum etc.)



# What are traditional hardware offload hints?

NIC hardware provides offload hints in RX (and TX) descriptors

- The netstack SKB packet data-struct stores+uses these

RX descriptors can e.g. provide:

- RX-checksum validation, RX-hash value, RX-timestamp
- RX-VLAN provides VLAN ID/tag non-inline

TX descriptors can e.g. ask hardware to perform actions:

- TX-checksum: Ask hardware to compute checksums on transmission
- TX-VLAN: Ask hardware to insert VLAN tag
- Advanced: TX-timestamp HW stores TX-time and feeds back on completion
- Advanced: TX-LaunchTime ask HW to send packet at specific time in future

# What are XDP-hints

XDP-hints dates back to NetDevConf Nov 2017 (by PJ Waskiewicz)

- Purpose: Let XDP access HW offload hints

Basic idea:

- Provide or extract (from descriptor) NIC hardware offload hints
- Store info in XDP metadata area (located before pkt header)

XDP metadata area avail since Sep 2017 (by Daniel Borkmann)

- Space is limited (currently 32 bytes)

Main reason XDP-hints work stalled

- No consensus on layout of XDP metadata
- BTF was not ready at that time

# XDP-hints layout defined via **BTF layout**

My **proposal**: Use BTF to define the layout of XDP metadata

- Each NIC driver can choose its own BTF layout
- Slightly **challenging requirement**:
  - NIC driver can **change layout per pkt** (e.g timestamp only in PTP pkts)

Open question: **Will BTF be a good fit for this use-case?**

Next slides: Explaining **BTF technical details**





# Introducing BTF – BPF Type Format

BTF compact Type Format (based on compiler's DWARF debug type info)

- Great [blogpost](#) by Andrii Nakryiko
  - 124MB of DWARF data compressed to 1.5MB compact BTF type data
- Suitable to be [included in Linux kernel image](#) by default
  - See file `/sys/kernel/btf/vmlinux` avail in most distro kernels
- Kernel's runtime data structures have become [self-describing via BTF](#)

```
# bpftool btf dump file /sys/kernel/btf/vmlinux format c
```

# More components: CO-RE + BTF + libbpf

Blogpost on BPF CO-RE (Compile Once – Run Everywhere) (Andrii Nakryiko)

- Explains how BTF is one piece of the puzzle
- BPF ELF object files are made portable across kernel versions via CO-RE
- LLVM compiler emits BTF relocations (for BPF code accessing struct fields)

BPF-prog (binary ELF object) loader libbpf combines pieces

- Tailor BPF-prog code to a particular running kernel
- Looks at BPF-prog recorded BTF type and relocation information
  - matches them to BTF information provided by running kernel
  - updates necessary offsets and other relocatable data
- Kernel struct can change layout, iff member name+size stays same

# Code-Example: Partial struct + runtime BTF-id

BPF-prog can define **partial struct** with few members

- libbpf matches + "removes" **triple-underscore** after **real struct name**
- **preserve\_access\_index** will be matched against kernel data-structure

```
struct sk_buff local {  
    u32 hash;  
} __attribute__((preserve_access_index));  
  
SEC("kprobe/udp_send_skb.isra.0")  
int BPF_KPROBE(udp_send_skb, struct sk_buff local *skb)  
{  
    u32 h; u32 btf_id;  
    BPF_CORE_READ_INTO(&h, skb, hash); /* skb->hash */  
    btf_id = bpf_core_type_id_kernel(struct sk_buff local);  
    bpf_printk("skb->hash=0x%x btf_id(skb)=%d", h, btf_id);  
}
```

Notice: Can get **btf\_id** for **sk\_buff** used by **running kernel**

# BTF type IDs and their usage

BTF system has **type IDs** to refer to each-other (in compressed format)

- Zero is not a valid BTF ID and numbering (usually) **starts from one**
  - Userspace can dump and see numbering via **bpftool btf dump file**

**Kernel's BTF** data files are located in **/sys/kernel/btf/** (modules since **v5.11**)

- Main file **vmlinux** contains every type **compiled into kernel**
- All **module files** **offset** ID numbering to start at **last vmlinux ID**
  - Allows module to reference vmlinux type IDs (for compression)

**Userspace** BPF-prog ELF-object files also contains **BTF sections**

- This is known as **local** BTF and numbering starts at one
- BPF-prog can query own local BTF id via: **bpf\_core\_type\_id\_local()**



# Back to XDP-hints

Back to XDP-hints and XDP metadata area

# XDP metadata requirements

XDP metadata area has some **properties**

- Grows “backwards” from where packets starts
- Must be 4 byte aligned
- Limited size (currently) 32 bytes

BPF-prog can expand/grow area via helper: **bpf\_xdp\_adjust\_meta**

- pkt-data pointers are invalidated after calling this
- **Verifier** requires **boundary checks** to access metadata area

Common gotcha: Compiler likes to pad C-struct ending

- Avoid/fix via: **\_\_attribute\_\_((packed))**



# Expected users of the XDP-hints

Users/consumers of XDP-hints in BTF layout

- BPF-progs first obvious consumer (either XDP or TC hooks)
- XDP to SKB conversion (in veth and cpumap) for traditional HW offloads
  - e.g. RX-hash, RX-checksum, VLAN, RX-timestamp
  - Can potentially simplify NIC drivers significantly
- Chained BPF-progs can communicate state via metadata
- AF\_XDP can consume BTF info in userspace to decode metadata area



# Motivation for XDP to SKB conversion

**Moonshot:** NIC drivers without SKB knowledge

- End-goal with XDP to SKB conversion
- Make it possible to write NIC drivers Ethernet L2 “only”

Pros: **Avoids** taking the SKB “**socket**” **overhead** at driver level

- Next step: Speedups Linux bridging and routing (with xdp\_frame)
  - Meaning: Normal Linux netstack get speedup for routing use-cases

# Hardware motivation and considerations

Goal: Hardware should produce XDP-hints

- Easy as DMA area next to metadata

Consider defining Endianness: Big vs Little endian

- In XDP-hints struct layout
- Given BTF is flexible, can be added later when HW appears

# XDP-hints exploring solutions using BTF

Design not fully done yet

- Upstream interaction will likely change solution anyhow

Next slides: Proposed solutions with pros and cons

# Solution(A): Internal kernel focus

Kernel-side: Could **extend** **xdp\_buff** + **xdp\_frame** with "btf\_id" or ptr

- **Pros**: Gives BPF-prog access to reading **ctx->btf\_id**
  - **Cons**: **AF\_XDP** cannot read this **ctx->btf\_id**
  - **Unknown**: Can **chained BPF-progs** **update/write** **ctx->btf\_id** ?
- **Pros**: **Kernel internally** can (likely) store **pointer to btf** struct
  - **Cons**: This needs reference counting and **race/lifetime handling**
- **Pros**: **XDP to SKB** conversion (should be) easier to extract offloads
  - **Unknown**: Will it be harder to support different **layout per pkt** ???

Details: If **dropping** requirement **layout per pkt**

- Possible to **store btf pointer in** **xdp\_rxq\_info**

# Solution(B): Decouple with btf\_id in metadata

Place "btf\_id" value inside metadata area, as last member

- last member: due to "grows" backwards, important for AF\_XDP decoding
- Extend xdp\_buff + xdp\_frame (+AF\_XDP) with flags that BTF is "enabled"
  - Notice: Need 3 flags for BTF "origin" (vmlinux, module or local)
    - module resolved via xdp\_rxq\_info or xdp\_frame->dev\_rx

This achieves decoupling via btf\_id as it becomes struct's "version" number

- Pros: Easy to handle different layout per pkt
  - as BPF-prog (or AF\_XDP) can multiplex on btf\_ids known to "them"
- Cons: XDP to SKB conversion harder as kernel cannot trust btf\_id
  - Solution: Add new BPF hook at XDP to SKB point (BPF prog builds skb)



# Solution(C): Combined proposal

Still place "btf\_id" value inside metadata area, as last member

- BUT is considered cached version from kernel stored pointer to btf
  - Cons: reference counting and lifetime handling still needed

Pros: XDP to SKB conversion can work

- via check if btf\_id matches btf-ptr id before trusting BTF layout
- Details: For "lifetime" module BTF driver could disable this step on teardown

Pros: Chained BPF-progs works

- (last\_member) btf\_id becomes a communication channel
  - Unknown: How to communicate BTF "origin" (vmlinux, module or local)?

# What BTF layout does a driver provide?

How to solve “exporting” available BTF-layouts

- per NIC driver

Is a new UAPI needed?

# What BTF layout does this driver provide?

How does userspace (and libbpf) know:

- What BTF layout does this driver provide?

New **UAPI** **might not** be needed:

- Remember: **BTF info** avail via **/sys/kernel/btf/**
  - both for **vmlinux** and **modules**
- libbpf parses and resolves relocations via these

Struct **naming-convention** for xdp\_hints

- Could be way for drivers to “export” available BTF-layouts?



# Proposal: Encapsulating C-code union?

Each NIC driver could have a `union` named `xdp_hints_union`

- Structs added to union, means driver `may` use this BTF layout
- Notice: Union "sub" structs automatically gets own BTF IDs
- Essentially: Way to describe/support NIC using `layouts per packet`

Complications: `metadata grows backwards`

- `Padding` needed if union should `match memory layout`
  - Cons: Union padding quickly gets "ugly" in C-code
  - Pros: Easier for driver C-code with one type for metadata area

# Define “generic” xdp\_hints common struct

Idea: **Partly UAPI/kABI** approach

**Kernel (not module) struct xdp\_hints\_common** (vmlinux BTF id)

- Should cover today’s **known SKB offload hints**
- Could have **some defines as UAPI** in bpf.h (e.g. flags hash-type, csum-type)

**NIC drivers** can let their **xdp\_hints** include **common struct as member**

- Can **extend** with **NIC specific hints** by adding **flag** in xdp\_buff/xdp\_frame
  - That indicates **layout is compatible** with “xdp\_hints\_common”

**Pros:** Easier to implement **XDP to SKB transition**

- **Cons:** Goes against the fully dynamic BTF based layout

# Metadata + BTF = communication channel

Using **metadata** area to **communication state**

- Create **structure via BTF**

Relevant for:

- **Chained BPF-progs** (between XDP to TC hooks, also BPF tail-calls)
- **XDP-prog** sending info to userspace **AF\_XDP** sockets

Already: **Works today!**





# Example use-case: XDP-prog to AF\_XDP

Code: How to transfer info from XDP-prog to AF\_XDP via BTF

- GitHub XDP-project: bpf-examples/ AF\_XDP-interaction
- Use-case: XDP RX-timestamp for Real-Time TTEthernet sync (PCF frames)
- Changes layout per packet via btf\_id as last member

Uses: local BTF info in BPF-prog ELF object

- Thus, no kernel extensions needed
- BPF-prog gets own local BTF id via: bpf\_core\_type\_id\_local()

Shows userspace C-code decoding BTF format

- Extracts offset + size for named struct members

# End: Questions?

## Resources:

- XDP-project - [GitHub.com/xdp-project](https://github.com/xdp-project)
  - Get an easy start with [xdp-project/bpf-examples](https://github.com/xdp-project/bpf-examples)
- XDP-hints mailing list: [xdp-hints @ xdp-project.net](mailto:xdp-hints@xdp-project.net)
  - <https://lists.xdp-project.net/>

# Extra slides

# Traditional hint: RX-hash – implementation details

Kernel SKB: Hash value (only) 32-bit

- pkt\_hash\_types: PKT\_HASH\_TYPE\_{NONE,L2,L3,L4} (bit skb->l4\_hash)
- SKB bit (skb->sw\_hash) if software computed hash

Hardware provides RSS-hash (standardised by Microsoft)

- Kernel drops info, RSS hashing type identify:
  - Tell us if this is IPv4: UDP or TCP
  - IPv6: UDP or TCP, and if extension headers are present

# Traditional hint: **VLAN** both RX and TX – **impl.**

**net\_device feature flags** for **enabling** VLANs offload hints:

- `NETIF_F_HW_VLAN_{CTAG,STAG}_{RX,TX}_BIT`
- C-tag = inner tag Customer-tag, S-tag = outer tag Service-provider-tag

Kernel **SKB fields**:

- `vlan_present` (1-bit), `vlan_tci` (16-bits), `vlan_proto` (BE 16-bit)

`ethtool -{show-}features: rx-vlan-offload + tx-vlan-offload`

# Traditional hint: RX-checksum – impl. details

net\_device feature flags avail for checksum capabilities

- e.g. NETIF\_F\_{HW,IP,IPV6}\_CSUM

Kernel stores checksum type (in `skb->ip_summed` 2-bits)

- CHECKSUM\_{NONE,UNNECESSARY,COMPLETE,PARTIAL}

CHECKSUM\_COMPLETE fills `skb->csum` (union with `csum_start+csum_offset`)

CHECKSUM\_PARTIAL needs `skb_checksum_start_offset` (`skb->csum_start`)

- Depends on `skb->csum_start + skb->csum_offset`

Encap: `skb->csum_level` consecutive checksums found in pkt

- Minus ones verified as CHECKSUM\_UNNECESSARY
- Encap case also sets bit `skb->encapsulation`



# Traditional hint: TX-checksum – impl. details

Requests net\_device to update packet checksum fields

- SKB reusing types in `skb->ip_summed` (2-bits)
- features: `NETIF_F_{IP,IPV6,HW}_CSUM`

CHECKSUM\_PARTIAL meaning at TX

- Do checksumming from `skb->csum_start` up to the end
- In pkt store checksum at offset `skb->csum_start + skb->csum_offset`