

XDP and BPF insights

Programmable Runtime Extending Linux Kernel for Packet Processing

Jesper Dangaard Brouer
Senior Principal Kernel Engineer

Tech Summit
October 2021

Overview: What will you learn?

What is BPF really ?

- How this technology fundamentally changes existing OS-model

Taming BPF superpower – is not easy

- BPF “user experience” – could be better!

What is XDP?

- and what pain points have recently been resolved

What is AF_XDP?

- How is this connected with XDP and deep-dive into tech details

What is BPF ?

From: <https://ebpf.io/what-is-ebpf>

eBPF is a revolutionary technology that can run sandboxed programs in the Linux kernel without changing kernel source code or loading a kernel module

BPF is a **technology name**: no longer an acronym

Rate of innovation at the operating system level: **Traditionally slow**

- BPF enables things at OS-level that were not possible before
- BPF will **radically increase** rate of innovation

Traditional Kernel development process


Application Developer:

I want this new feature to observe my app



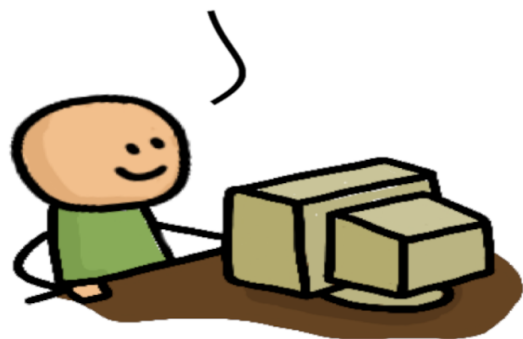
Hey kernel developer! Please add this new feature to the Linux kernel

OK! Just give me a year to convince the entire community that this is good for everyone.



1 year later...

I'm done. The upstream kernel now supports this.



But I need this in my Linux distro



5 year later...

Good news. Our Linux distribution now ships a kernel with your required feature

OK but my requirements have changed since...



BPF development process

Application Developer:

I want this new feature to observe my app



eBPF Developer:

OK! The kernel can't do this so let me quickly solve this with eBPF.



A couple of days later...

Here is a release of our eBPF project that has this feature now. BTW, you don't have to reboot your machine.



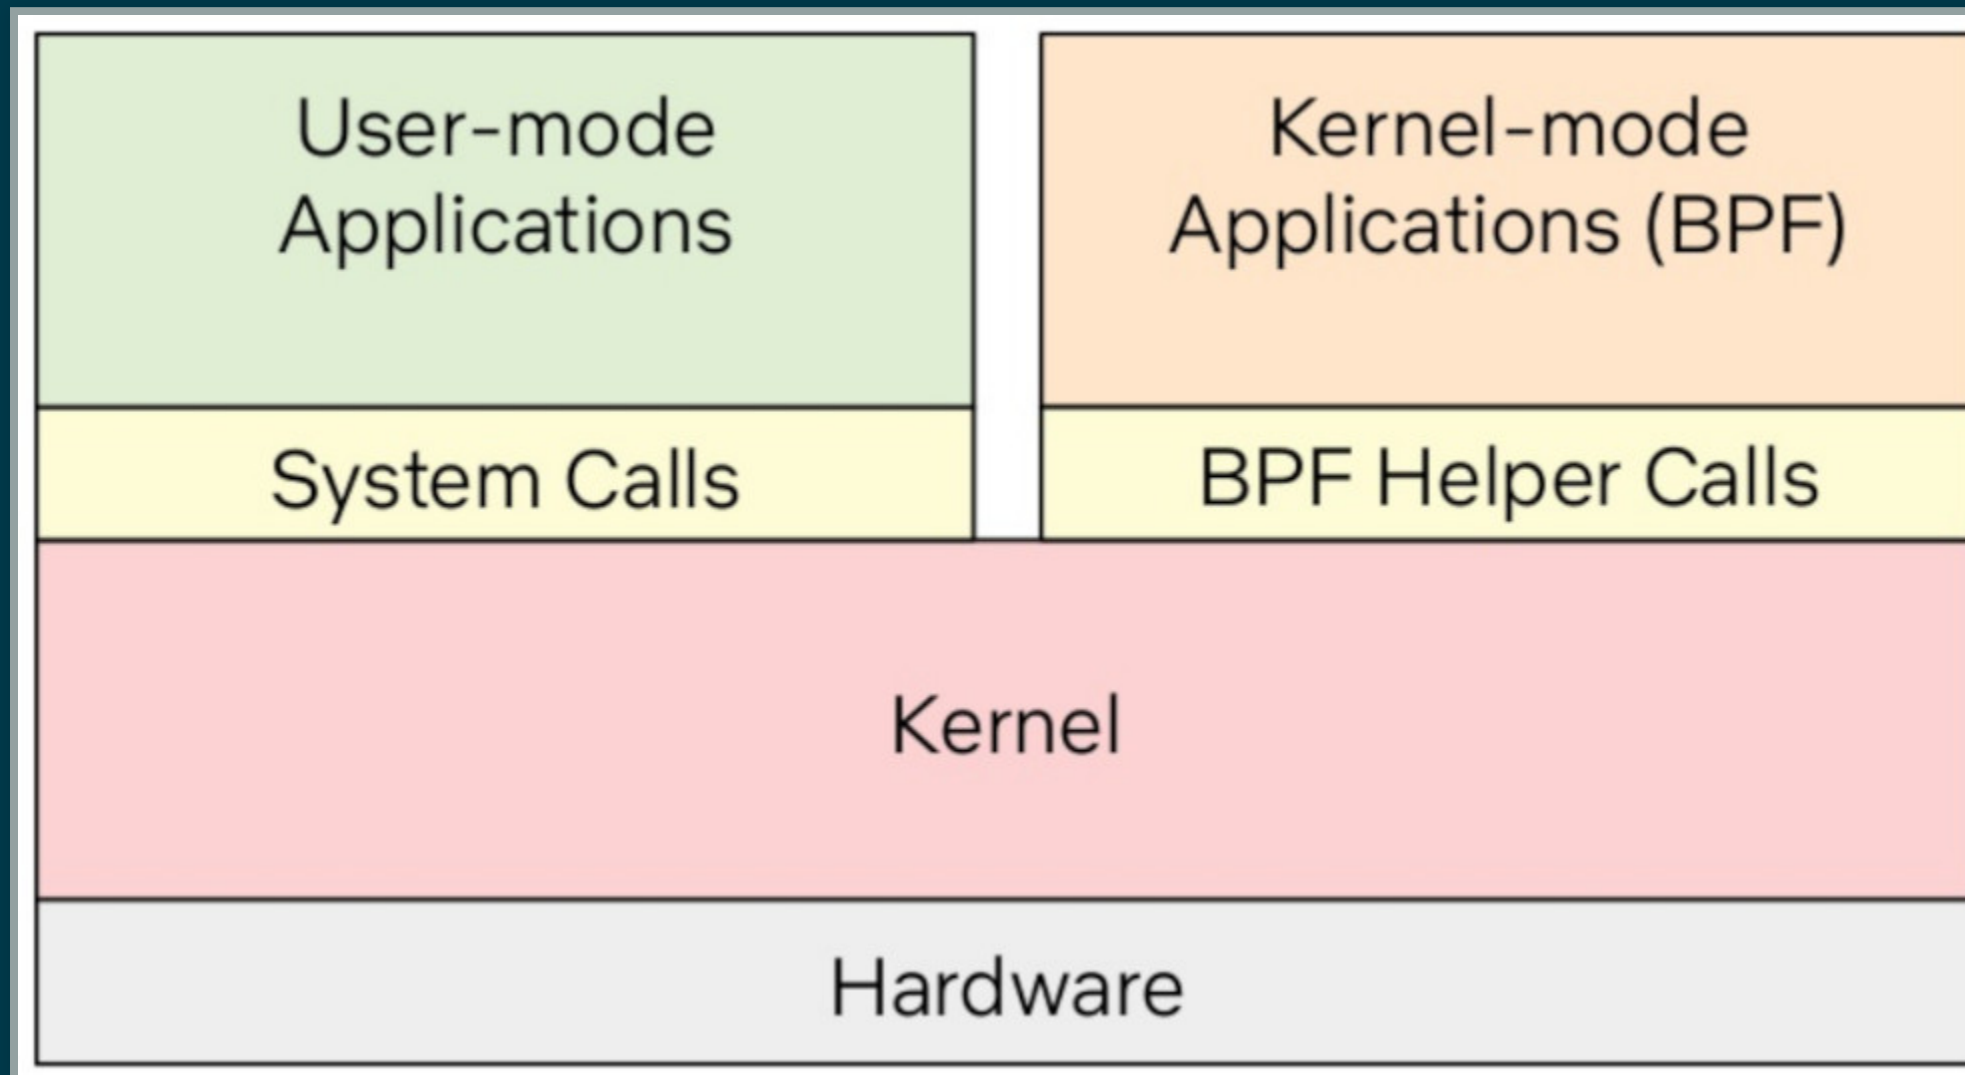
What is BPF from OS-vendor perspective

From an OS vendor (like Red Hat) perspective

- BPF is a **fundamental change** to a 50-year old **kernel model**
 - New interface for applications to make kernel requests, alongside syscalls
- BPF is “**kernel-mode applications**”
- See it as **safer** alternative to **kernel modules**
 - Better way to extend the kernel dynamically
 - Kernel developers: Do lose some control
 - ... but positive for rate of innovation

A new Operating System Model

Modern Linux: becoming Microkernel-ish (from Brendan Gregg)



Adjust mental model of system development

As System Developers: Adjust your mental model

- No need to bring everything into userspace
- Can task be solved via kernel-side inline processing and manipulation?

Utilize BPF superpowers

- Linux Kernel has become a flexible tool at your disposal
- Run your own BPF kernel-mode application inside the kernel

BPF components

Closer look at the BPF components:

- **Bytecode** – Architecture independent **Instruction Set**
 - **JIT** to native machine instructions (after loading into kernel)
- **Runtime environment** – Linux kernel
 - **Event based** BPF-hooks all over the kernel
 - Per hook limited access to kernel functions via **BPF-helpers**
- **Sandboxed** by the BPF **Verifier**
 - Limits and verifies memory access and instructions limit

BPF **concepts**: context, maps and helpers

Each BPF **runtime hook** gets a **pointer to a context** struct

- BPF bytecode has access to context (read/write limited)
 - Verifier adjusts bytecode when accessing struct members (for safety)

The BPF program itself is stateless

- BPF **maps** can be used to **create state**
- Maps are basically **key = value** containers
- Maps can hide complex features implemented on Kernel side

BPF helpers are used for

- **Calling Kernel functions**, to obtain info/state from kernel

BPF functional areas

BPF has hooks ALL over the kernel

- More interesting: What **functional areas** are these being used?

What (current) areas are BPF being used within?

- **Networking** (ahead as BPF started in networking)
 - Network control and data plane implemented in BPF (TC-BPF/XDP)
 - TCP Congestion Control implemented in BPF
- **Security**
- **Observability** (tracing)

Happy 7th Birthday BPF

Happy birthday BPF!

- 7 years old (See Alexei's post [Sunday 26 September 2021](#))
- XDP initial commit is approx 5 years + 2 months

Exciting things ahead

- [eBPF Foundation](#) (ebpf.io/charter) working towards **standardisation**
- Microsoft Windows introduce BPF in their kernel

Digital age and out-of-date documentation

BPF + LLVM features evolved over time

- Google search results: Many but out-dated articles
- See outdated approaches are used as best-practices :-(

Some quick advice to follow

- Use latest LLVM compiler (and -target bpf)
- Install latest pahole tool (used for BTF generation)
- Get Kernel with BTF (BPF Type Format) support
- Use new BPF-maps definitions (".maps" section) with BTF support

Taming **BPF superpowers** – not easy

BPF superpowers – not easy to use – sorry

- Gain kernel level building block, that can be safely updated runtime
- Taming and learning-curve is challenging

BPF makes extending Kernel **easier** than Kernel modules

- Don't confuse this with "easy to use"

BPF development is hard

Know this: BPF development is hard

- **Mental model** mind-shift: Coding “kernel-mode applications”
- Requires understanding internal kernel functions
- Available BPF **features** depend on LLVM **compiler versions**
- Developers will experience: **Verifier rejecting** programs
- Coding in Restricted C and passing **verifier is frustrating**
 - Corner-cases due to LLVM can be the issue
- Troubleshooting event based BPF-prog running kernel side is challenging

BPF user experience – could be better

BPF is great revolutionary technology!

- BUT end-user deployment experience can be rough

Recommend watching recent LPC 2021 talk by CloudFlare

- Talk: BPF user experience rough edges
- Covers 9 common pitfalls (with sad pandas)

BPF communities

Remember to **reach out to BPF communities** when stuck

- BPF Kernel developers: <mailto:bpf@vger.kernel.org>
- Slack channel: <https://ebpf.io/slack>
- LLVM compiler questions: <mailto:iovisor-dev@lists.iovisor.org>

XDP communities

- XDP-newbies: <mailto:xdp-newbies@vger.kernel.org>
- GitHub project: <https://github.com/xdp-project>
- IRC on oftc.net channel **#xdp**

BPF example code

Best documentation is BPF example code

Under XDP-project: github.com/xdp-project/

- [bpf-examples](#) - Practical BPF examples and build environment
- [xdp-tutorial](#) - Tutorial with assignments (**Warning**: uses old BPF-maps)
- [xdp-tools](#) - Tools (xdpdump) + libxdp for **multiple XDP-progs** on interface
- [xdp-cpumap-tc](#) - Show XDP + TC-BPF solving Qdisc lock scaling

BPF networking

Focus on BPF for networking

- **XDP** (eXpress Data Path) is **our focus**
- **TC-BPF** hooks are **equally important** for practical use-cases
- BPF hooks for cgroups can also be useful for containers

Why was an eXpress Data Path (XDP) needed?

Linux **networking stack** assumes layers **L4-L7** are needed for every packet

- Root-cause of slowdown: (relative) high initial RX cost per packet

Needed to stay relevant as NIC speeds increase (time between packet small)

- New faster and earlier networking layer was needed to keep up.

XDP operate at layers **L2-L3**

- **L4 load-balancer** possible when **no IP-fragmentation** occurs

If you forgot OSI model:

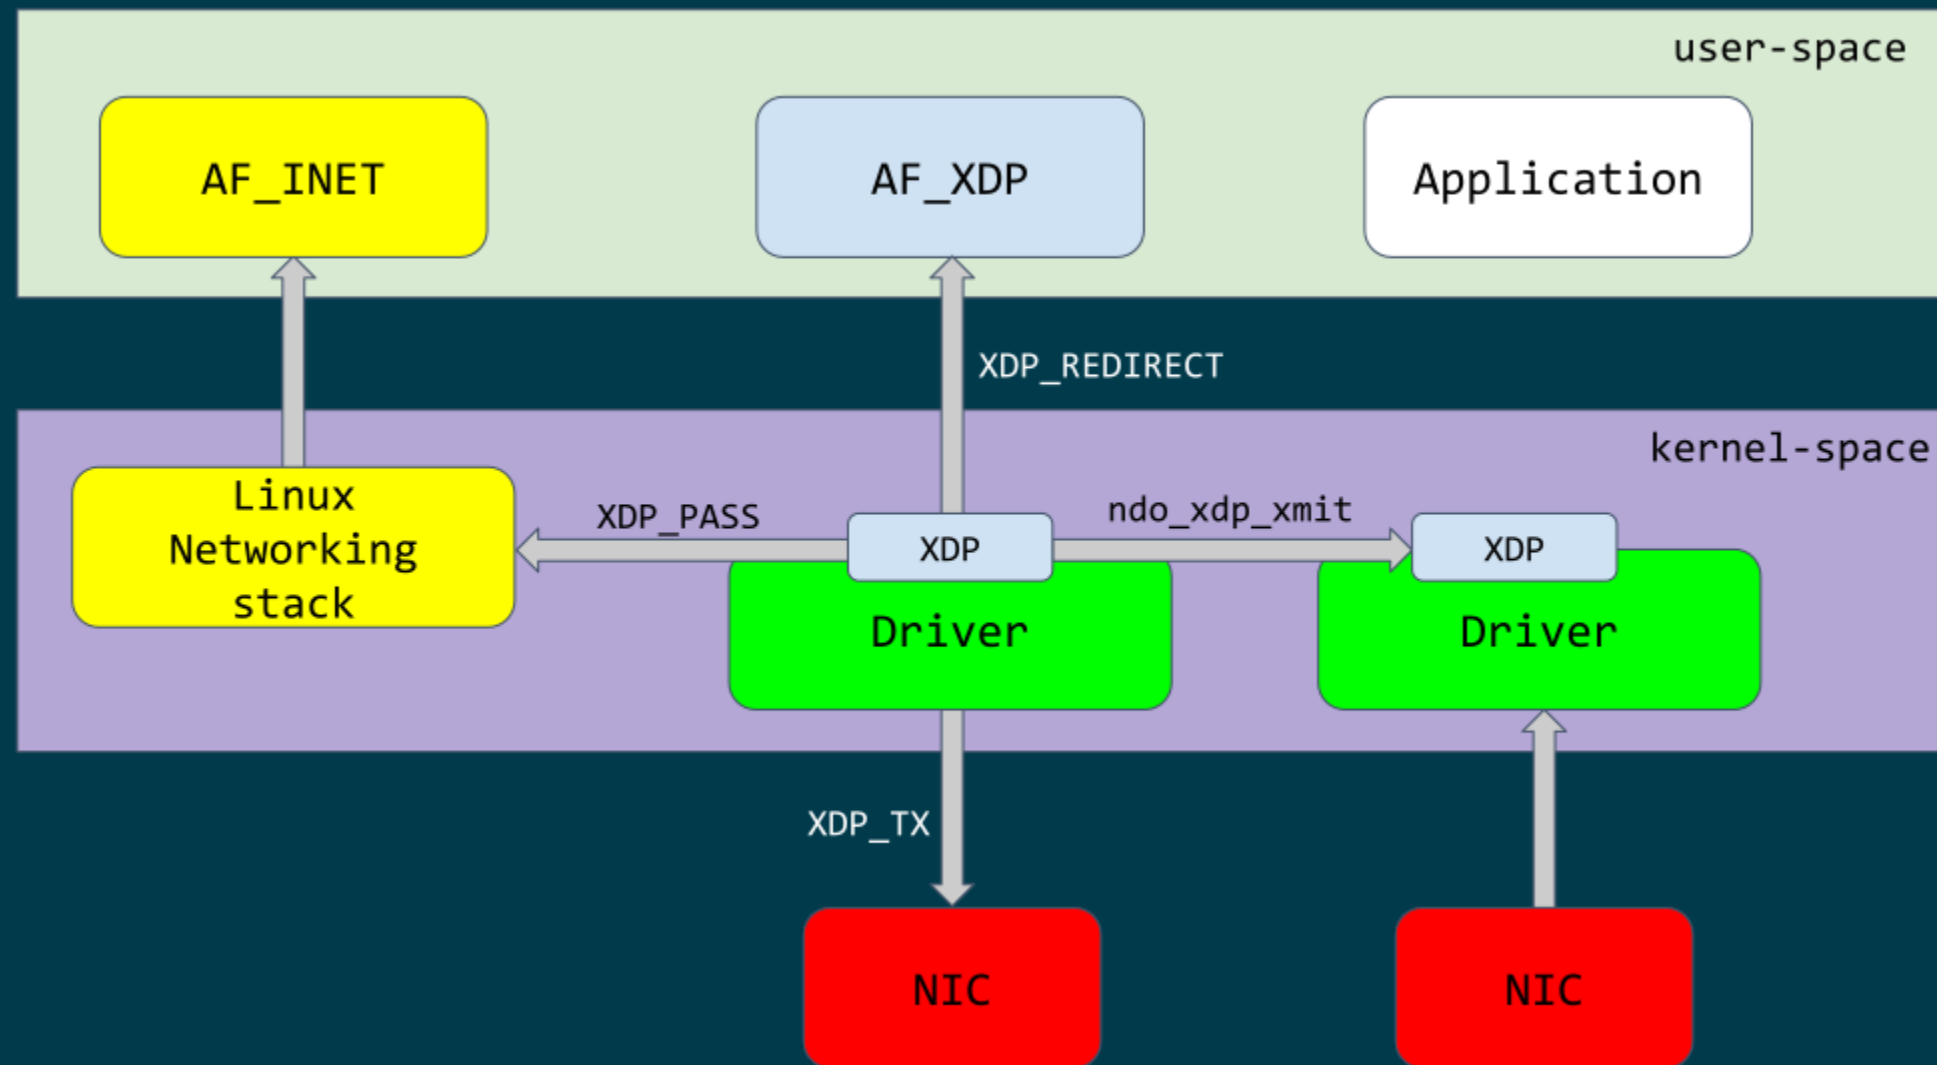
- L2=Ethernet
- L3=IPv4/IPv6
- L4=TCP/UDP
- L7=Applications

What is XDP?

XDP (eXpress Data Path) is a Linux **in-kernel** fast-path

- **New programmable layer in-front** of traditional network stack
 - Read, modify, drop, redirect or pass
 - For L2-L3 use-cases: seeing **x10 performance** improvements!
- **Avoiding memory allocations**
 - No SKB allocations and no-init (SKB zeroes 4 cache-lines per pkt)
- Adaptive **bulk** processing of frames
- Very **early access** to frame (in driver code **after DMA sync**)
- Ability to **skip (large parts) of kernel code**
 - Evolve XDP via **BPF-helpers**

XDP architecture




Performance graphs

Performance measurements taken from our [XDP-paper](#)

System used for testing

- Intel(R) Xeon(R) CPU **E5-1650** v4 @ **3.60GHz**
- NIC driver **mlx5**: Mellanox **ConnectX-5** Ex (MT28800)

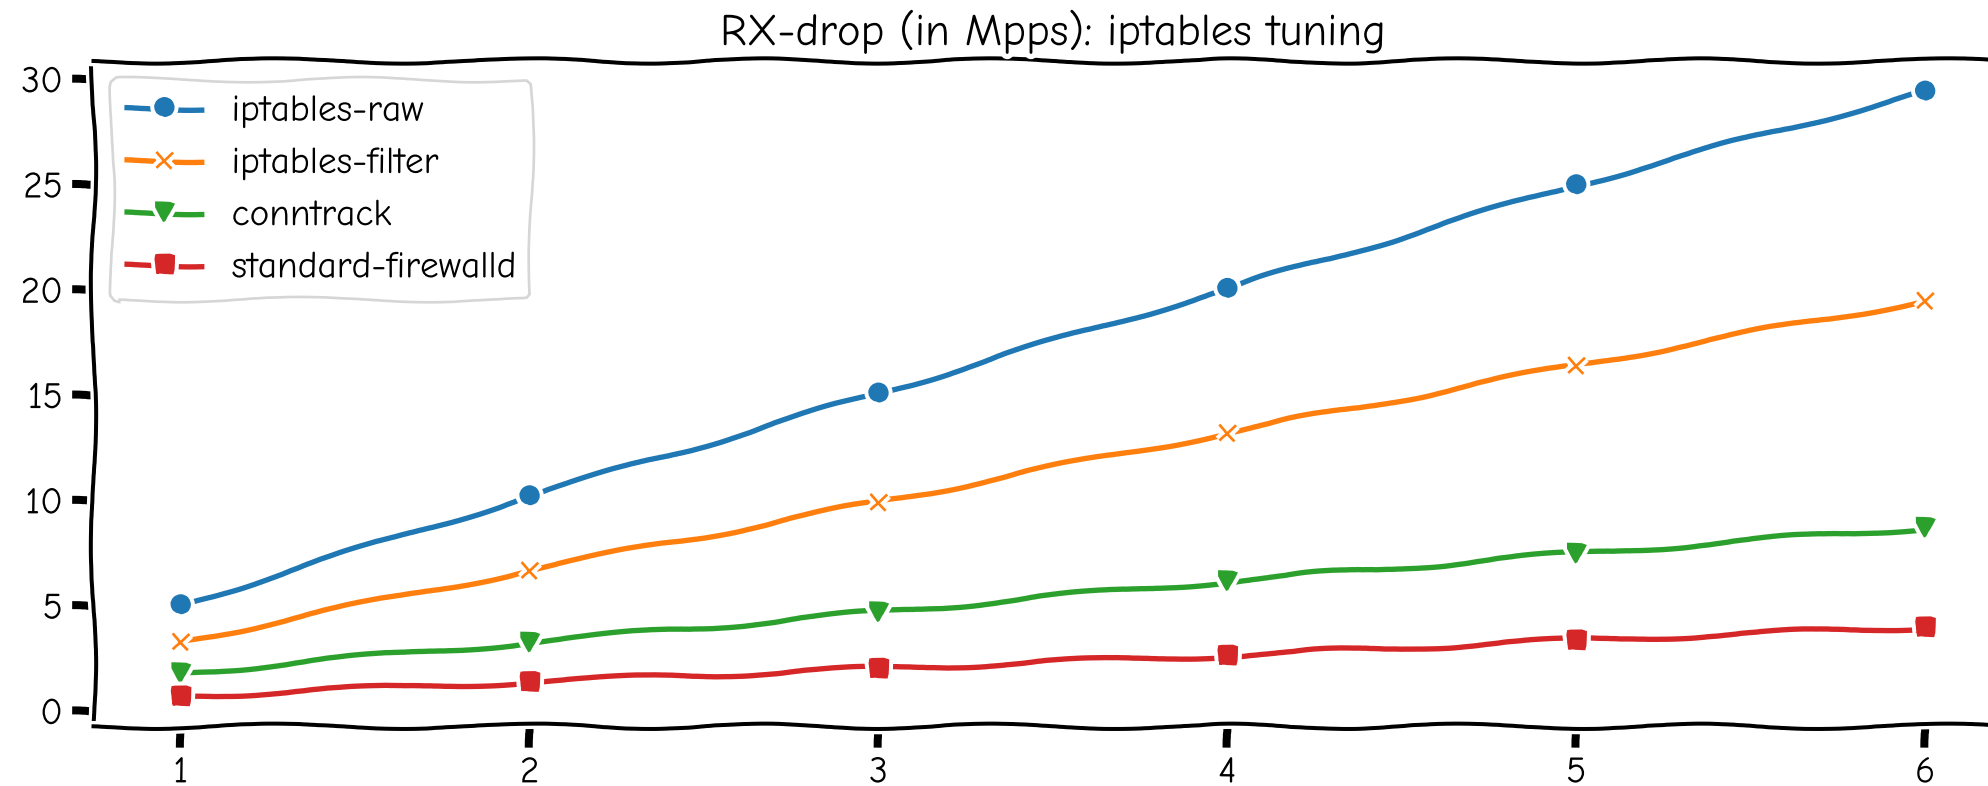
XDP performance



XDP_DROP: 100Gbit/s mlx5 max out at **108 Mpps** (CPU E5-1650v4 @3.60GHz)

- **PCIe tuning needed** – NIC compress RX-descriptors (rx_cqe_compress on)

Zoom-in: on iptables performance tuning



iptables can be tuned to perform and scale well

- Especially if avoiding involving conntrack

XDP performance: both latency and throughput

XDP throughput and packet-per-second (PPS) performance super

- Real design goal is improving latency and throughput at same time

Designed with adaptive bulking

- Run as part of NAPI-poll (softirq) processing (max 64 frame budget)
- Pickup frames from RX-ring if available (no waiting for bulks)
- End of drivers NAPI-poll function flush any pending xdp_frames
- Another driver egress REDIRECT (devmap) flush every 16 frames
- REDIRECT to another CPU (cpumap) flush every 8 frames
- As load increase, bulking opportunities happen, system scale to load

XDP scaling across CPUs

XDP **redirect** to **another CPU** (via BPF **cpumap** type)

- Scalability mechanism: let XDP control on that CPU netstack runs

Allow to combine **fast DDoS** handing and **slower netstack** on **same hardware**

- Some CPUs run dedicated **fast-path** packet processing
- Delegate to other CPUs via XDP-redirect
- **Remote CPUs** receive **raw xdp_frames**, next steps:
 - Can (optionally) run XDP-prog for further filtering
 - **Bulk allocate SKBs** and **start network stack** on this **CPU**
- **Attacker** hitting netstack/app **slow-path** cannot influence fast-path

What is AF_XDP?

What is **AF_XDP**? (the **Address Family XDP** socket)

- Hybrid **kernel-bypass** facility, **selectively move frames out of kernel**
- XDP/BPF-prog filters packets using **REDIRECT** into AF_XDP socket
- Delivers raw L2 frames into userspace (via memory mapped ring buffer)

Realize: **in-kernel** XDP **BPF-prog step** opens **opportunities**

- Can augment/modify packets prior to AF_XDP delivery
 - E.g. record a **timestamp** at this **early** stage
- Use CPUMAP redirect: Move **netstack traffic** to **other CPUs**

WARNING: Next slides: **Deep dive** into AF_XDP **details**

- Most casual readers of slide deck can **skip these details**

Where does AF_XDP performance come from?

Lock-free channel directly from driver RX-queue into AF_XDP socket

- Single-Producer/Single-Consumer (SPSC) descriptor ring queues
- Single-Producer (SP) via bind to specific RX-queue id
 - NAPI-softirq assures only 1-CPU process 1-RX-queue id (per sched)
- Single-Consumer (SC) via 1-Application
- Bounded buffer pool (UMEM) allocated by userspace (register with kernel)
 - Descriptor(s) in ring(s) point into UMEM
 - No memory allocation, but return frames to UMEM in timely manner
- Transport signature Van Jacobson talked about
 - Replaced by XDP/eBPF program choosing to XDP_REDIRECT

Details: Actually **four** SPSC ring queues

AF_XDP **socket**: Has **two rings**: **RX** and **TX**

- Descriptor(s) in ring points into UMEM

UMEM consists of a number of equally sized chunks

- Has **two rings**: **FILL** ring and **COMPLETION** ring
- FILL ring: application gives kernel area to RX fill
- COMPLETION ring: kernel tells app TX is done for area (can be reused)

Gotcha by RX-queue id binding

AF_XDP bound to **single RX-queue id** (for SPSC performance reasons)

- NIC by default spreads flows with RSS-hashing over RX-queues
 - Traffic likely not hitting queue you expect
- You **MUST** configure NIC **HW filters** to **steer to RX-queue id**
 - Out of scope for XDP setup
 - Use ethtool or TC HW offloading for filter setup
- **Alternative** work-around
 - Create as many **AF_XDP** sockets as RXQs
 - Have userspace poll()/select on all sockets

XDP pain points **resolved**

Followup to Linux Plumber 2019: **XDP** the distro view

- Some of the **pain points** have been **resolved**

Multiple XDP programs on a single interface

Followup to Linux Plumber 2019: [XDP the distro view](#)

The library [libxdp](#) (available via [xdp-tools](#))

- Have option of loading **multiple XDP programs** on a **single interface**
- See [dispatcher](#) API (`xdp_multiprog__*`) in README
- Depend on kernel feature `freplace` (read as: function replace)

XDP “tcpdump” packet capture

Tool ‘tcpdump’ does **not see all packets** anymore.

- E.g XDP_DROP and XDP_REDIRECT etc.

New tool ‘**xdpdump**’ (available via **xdp-tools**)

- Debug XDP programs already loaded on an interface
- Packets can be **dumped/inspected**:
 - **Before** on **entry** to XDP program
 - **After** at **exit** from an XDP program
 - Furthermore: at **exit** the XDP **action** is also **captured**
 - **Can inspect XDP_DROP packets!**
- Use Kernel features `fentry` + `fexit`
 - Also works with multi-prog dispatcher API

XDP future development

XDP multi-buff

- Allowing larger MTUs, Jumbo-frames and GRO/GSO compatibility

XDP-hints

- Extracting NIC hardware hints (from RX-descriptor)
- Traditional hints: RX-hash, RX-checksum, VLAN, RX-timestamp

Drivers without SKB knowledge

- based only on `xdp_frame`
- Depend on both XDP-hints + XDP multi-buff

End: Questions?

Resources:

- XDP-project - [GitHub.com/xdp-project](https://github.com/xdp-project)
 - Get an easy start with [xdp-project/bpf-examples](https://github.com/xdp-project/bpf-examples)