

Always-on latency monitoring with eBPF

Toke Høiland-Jørgensen
Red Hat

Acknowledgements

This is joint work with Simon Sundberg, Anna Brunstrom, Simone Ferlin-Reiter, Jesper Dangaard Brouer and Robert Chacón.


What is eBPF ?



- *Bytecode* - Architecture independent **Instruction Set**
 - *JIT* to native machine instructions (after loading into kernel)
- *Runtime environment* - Linux kernel
 - **Event based** BPF hooks all over the kernel
 - Per hook limited access to kernel functions via **helpers and kfuncs**
- *Sandboxed* by the eBPF *verifier*
 - Limits and verifies memory access and instructions limit

Passive latency monitoring

Kathie Nichols' [pping](#) + eBPF = [ePPing](#)



Sundberg, S. et al.: [Efficient Continuous Latency Monitoring with eBPF](#). Passive and Active Measurement (PAM 2023).

ePPing in action

```
$ sudo ./pping -i eth0
Starting ePPing in standard mode tracking TCP on eth0
15:02:47.835948282 TCP 198.49.23.145:443+45.145.92.2:58188 opening due to first observed packet from dest
15:02:47.859153254 TCP 45.145.92.2:58188+198.49.23.145:443 opening due to first observed packet from dest
15:02:47.885873388 21.2851 ms 21.2851 ms TCP 198.49.23.145:443+45.145.92.2:58188
15:02:50.248235439 21.1454 ms 21.1454 ms TCP 198.49.23.145:443+45.145.92.2:58188
^C
```

Cool! But doesn't really scale so well...

Enabling *always on* monitoring

```
$ sudo ./pping -i eth0 --aggregate 10
Starting ePPing in standard mode tracking TCP on eth0
Aggregating RTTs in histograms with 250 4 ms wide bins every 10 seconds
15:10:45.084560144: 198.49.23.0/24 -> rxpkts=53, rxbytes=102917, txpkts=51, txbytes=5233, rtt-count=4,
                                min=21.1356 ms, mean=25 ms, median=24 ms, p95=29.4 ms, max=29.6763 ms
15:10:45.084560144: 45.145.92.0/24 -> rxpkts=113, rxbytes=71232, txpkts=114, txbytes=116857
15:10:45.084560144: 0.0.0.0/0 -> rxpkts=26, rxbytes=4456, txpkts=24, txbytes=4485
15:10:45.084560144: ::/0 -> rxpkts=19, rxbytes=3215, txpkts=19, txbytes=3215
15:10:45.118005597: TCP=(pkts=202, bytes=181238), UDP=(pkts=31, bytes=8720), ICMP=(pkts=22, bytes=2156),
                    ECN=(Not-ECT=196, ECT1=49, ECT0=10)
^C
```

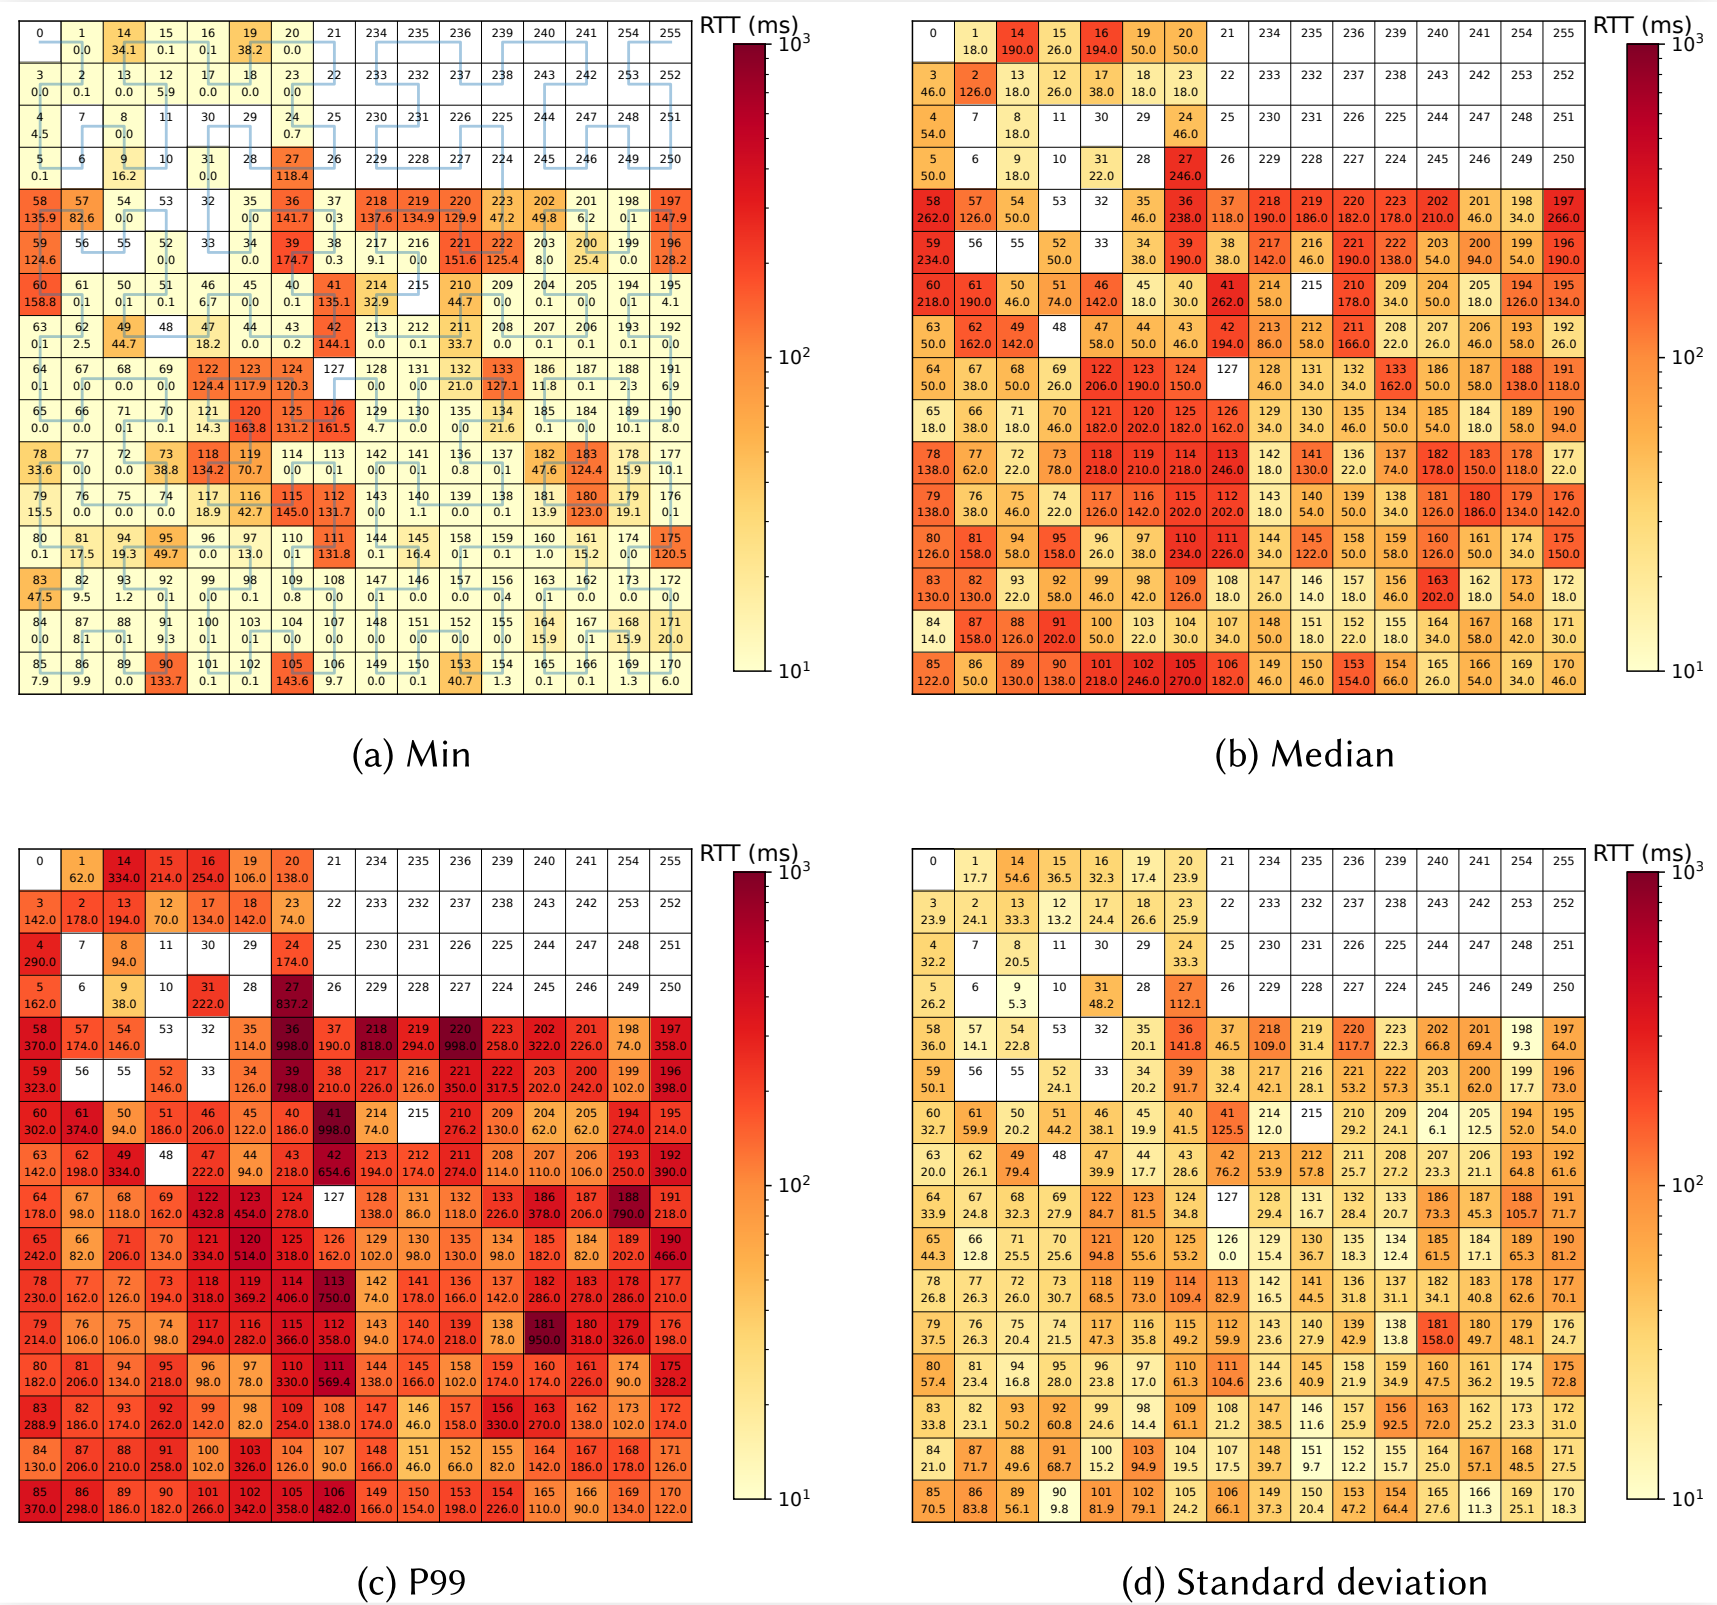
Histogram aggregation happens inside the kernel!

So what can we show with this?


One-month measurement study from a WISP middlebox (running [LibreQoS](#))

Under submission


Subnet RTTs



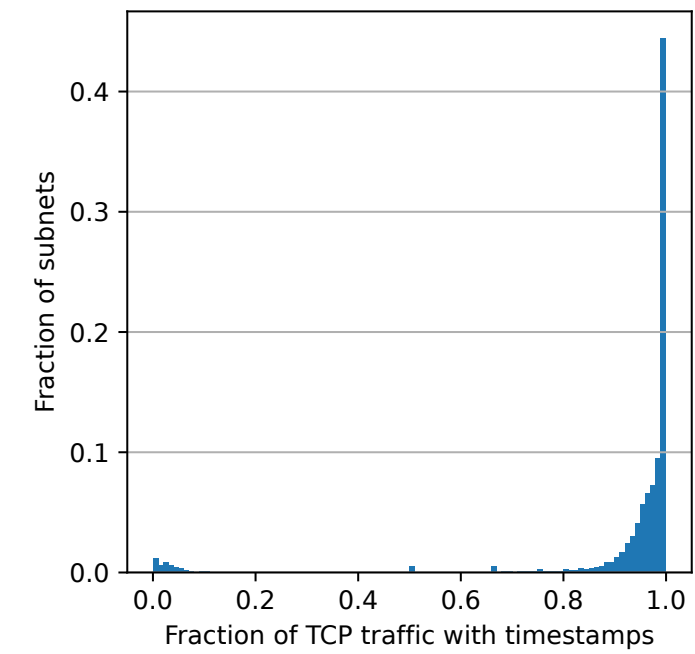
Traffic features



(a) Protocol occurrence




(b) ECN occurrence




(c) TCP timestamps

LAN and WAN traffic



(a) LAN



(b) WAN

Takeaways

With ePPing we can:

- Passively measure RTTs at scale on any Linux machine (server or middlebox)
- Aggregate metrics per subnet over the whole internet
- Investigate traffic characteristics on both the WAN and LAN side

What can this tell you about **your** network?

<https://github.com/xdp-project/bpf-examples/tree/master/pping>