

Achieving lower latency with eBPF and XDP

Toke Høiland-Jørgensen
Principal Kernel Engineer, Red Hat

Understanding Latency webinar
March 7th, 2023

What is eBPF ?

From: <https://ebpf.io/what-is-ebpf>

eBPF is a revolutionary technology that can run sandboxed programs in the Linux kernel without changing kernel source code or loading a kernel module

Rate of innovation at the operating system level: **Traditionally slow**

- eBPF enables things at the OS level that were not possible before
- eBPF can **radically increase** rate of innovation

eBPF components

Closer look at the eBPF components:

- **Bytecode** – Architecture independent **Instruction Set**
 - **JIT** to native machine instructions (after loading into kernel)
- **Runtime environment** – Linux kernel
 - **Event based** BPF hooks all over the kernel
 - Per hook limited access to kernel functions via **helpers and kfuncs**
- **Sandboxed** by the eBPF **verifier**
 - Limits and verifies memory access and instructions limit

What is XDP?

XDP (eXpress Data Path) is a Linux **in-kernel** fast path

- **Programmable layer in front** of the kernel networking stack
 - Read, modify, drop, redirect or pass
 - For L2-L3 use cases: seeing **x10 performance** improvements!
- **Avoiding memory allocations**
- Adaptive **bulk** processing of frames
- Very **early access** to frame (in driver code **after DMA sync**)
- Ability to **skip** (large parts) of kernel **code**

XDP performance

XDP_DROP: 100Gbit/s mlx5 max out at **108 Mpps** (CPU E5-1650v4 @3.60GHz)

Lower latency with eBPF

eBPF can help improve latency in a number of ways

- Lower latency by increasing PPS performance
- Custom software RSS steering
- Passive latency monitoring
- Future: Queueing for XDP



Scaling a latency-reducing traffic shaper

Use case: ISP middlebox providing per-customer bandwidth enforcement and bufferbloat mitigation (using kernel queueing infrastructure)

Problem: Software shaping doesn't scale because of **global qdisc lock**

Solution: XDP can choose which CPU to start the Linux networking stack on – steer a subset of customers to each CPU, so CPUs can run independently (avoiding the lock contention)

<https://github.com/xdp-project/xdp-cpumap-tc>

<https://github.com/LibreQoE/LibreQoS>



Passive latency monitoring

Use case: Monitor TCP traffic and extract flow latency (using TCP timestamps) to passively monitor traffic flowing through a middlebox.

Problem: The existing solution in software (**pping**) doesn't scale to high bandwidths

Solution: eBPF can inspect every packet with very low overhead – implement the monitoring in the kernel with eBPF, only export metrics to userspace

<https://github.com/xdp-project/bpf-examples/tree/master/pping>

Lower latency by increasing PPS performance

Use case: Linux-based servers exposed to the internet with high-speed NICs.

Problem: Head-of-Line (HOL) blocking in NIC RX rings causes latency and packet drops if kernel can't keep up with flow rate.

Solution: Implement filtering and redirection in XDP, allowing the kernel to keep up with the incoming packet rate.

Ex: <https://blog.cloudflare.com/l4drop-xdp-ebpf-based-ddos-mitigations/>

Future: Queueing for XDP

Use case: High-performance forwarding path with XDP “software offload”

Problem: XDP currently has no support for packet queueing and scheduling

Solution: We’re working on adding programmable queueing support to XDP

<https://lpc.events/event/16/contributions/1351/>



Closing remarks

eBPF allows **unprecedented visibility** into the OS, and **safe, dynamic extensibility** of core OS features, networking in particular.

eBPF unlocks the kernel's potential for innovation

- Pioneered on Linux, but exists in Windows too:
<https://github.com/microsoft/ebpf-for-windows>
- The eBPF Foundation (working on standardisation): <https://ebpf.foundation/>
- More examples of applications using eBPF: <https://ebpf.io/applications>
- Code examples: <https://github.com/xdp-project/bpf-examples>
- XDP tutorial: <https://github.com/xdp-project/xdp-tutorial>

End: Questions?

