

Introduction to xdp-tools

Toke Høiland-Jørgensen

Principal Kernel Engineer, Red Hat

NS team tech talk

June 8th, 2022

Reminder: What is XDP?

XDP (eXpress Data Path) is a Linux **in-kernel** fast-path

- **Programmable layer in-front** of traditional network stack
 - Read, modify, drop, redirect or pass
 - For L2-L3 use-cases: seeing **x10 performance** improvements!
- **Avoiding memory allocations**
 - No SKB allocations and no-init (SKB zeroes 4 cache-lines per pkt)
- Adaptive **bulk** processing of frames
- Very **early access** to frame (in driver code **after DMA sync**)
- Ability to **skip (large parts) of kernel code**
 - Evolve XDP via **BPF-helpers**

What is xdp-tools?

Upstream repo: <https://github.com/xdp-project/xdp-tools>

Contains:

- libxdp
- xdp-loader
- xdp-filter
- xdpdump
- xdp-traffigen (WiP)

Packaged in RHEL 8 & 9

libxdp and the multi-prog dispatcher

libxdp

libxdp is a **small library** built on top of libbpf; it provides:

- **User space support** code for using AF_XDP (moved from libbpf)
- **Multi-prog dispatcher** for running multiple XDP programs on a single interface

The multi-prog dispatcher

The kernel only allows executing a single XDP program per interface!

```
static volatile const struct xdp_dispatcher_config conf = {};
__attribute__((noinline)) int prog0(struct xdp_md *ctx) {
    volatile int ret = XDP_DISPATCHER_RETVAL;

    return ret; /* this function is replaced on load using freplace */
}
SEC("xdp")
int xdp_dispatcher(struct xdp_md *ctx)
{
    __u8 num_progs_enabled = conf.num_progs_enabled;
    int ret;

    if (num_progs_enabled < 1)
        goto out;
    ret = prog0(ctx); /* repeat for up to 10 progs */
    if (!(1U << ret) & conf.chain_call_actions[0])
        return ret;

out:
    return XDP_PASS;
}
```

The tools in xdp-tools

xdp-loader

General-purpose **program loader** based on libxdp (multi-prog support).

```
# xdp-loader load testns xdp_drop.o
# xdp-loader load testns xdp_pass_kern.o
# xdp-loader status
```

CURRENT XDP PROGRAM STATUS:

Interface	Prio	Program name	Mode	ID	Tag	Chain actions
lo		<No XDP program loaded!>				
wg0		<No XDP program loaded!>				
redhat0		<No XDP program loaded!>				
eth0		<No XDP program loaded!>				
testns		xdp_dispatcher	native	357	94d5f00c20184d17	
=>	50	xdp_drop		352	57cd311f2e27366b	XDP_PASS
=>	50	xdp_prog_simple		364	3b185187f1855c4c	XDP_PASS

xdp-filter

Simplistic **packet filter** demo (not a full-fledged firewall!)

```
# xdp-filter load testns -f ipv6
# xdp-filter ip fc00:dead:cafe:1::1
# xdp-filter status
```

CURRENT XDP-FILTER STATUS:

Aggregate per-action statistics:

XDP_ABORTED	0 pkts	0 KiB
XDP_DROP	5 pkts	0 KiB
XDP_PASS	3 pkts	0 KiB

Loaded on interfaces:

	Enabled features
xdpfilt_alw_ip	
testns (native mode)	ipv6,ipv4,allow

Filtered IP addresses:

	Mode	Hit counter
fc00:dead:cafe:1::1	dst	5

xdpdump

tcpdump-like utility, but **attaches to existing XDP programs**:

```
# xdpdump -i testns --rx-capture entry,exit
listening on testns, ingress XDP program ID 400 func xdpfilt_alw_ip, capture mode entry/exit, capture size 262144
1654692789.575086848: xdpfilt_alw_ip()@entry: packet size 118 bytes on if_index 19, rx queue 0, id 1
1654692789.575092034: xdpfilt_alw_ip()@exit[DROP]: packet size 118 bytes on if_index 19, rx queue 0, id 1
1654692790.590477081: xdpfilt_alw_ip()@entry: packet size 118 bytes on if_index 19, rx queue 0, id 2
1654692790.590493371: xdpfilt_alw_ip()@exit[DROP]: packet size 118 bytes on if_index 19, rx queue 0, id 2

# xdpdump -i testns -w - | tcpdump -nr -
listening on testns, ingress XDP program ID 400 func xdpfilt_alw_ip, capture mode entry, capture size 262144 bytes
reading from file -, link-type EN10MB (Ethernet), snapshot length 262144
14:54:06.190611 IP6 fc00:dead:cafe:1::2 > fc00:dead:cafe:1::1: ICMP6, echo request, id 59320, seq 1, length 64
14:54:07.216884 IP6 fc00:dead:cafe:1::2 > fc00:dead:cafe:1::1: ICMP6, echo request, id 59320, seq 2, length 64
```

These packets **were filtered**, but we can **still see them**!

xdp-trafficgen

XDP-based **programmable traffic generator** (not in RHEL yet):

```
# xdp-trafficgen udp ens3f1 # single core
[..]
XDP_REDIRECT      11150720 pkts (   8919659 pps)      696920 KiB (   4567 Mbits/s)

# xdp-trafficgen udp ens3f1 -t 6 # 6 cores
[..]
XDP_REDIRECT      65123603 pkts (  52095122 pps)     4070225 KiB ( 26673 Mbits/s)

# xdp-trafficgen udp ens3f1 -t 6 -d 100 # spraying over 100 dst-ports
[..]
XDP_REDIRECT      8226576 pkts (  32896120 pps)      514161 KiB ( 16843 Mbits/s)

# xdp-trafficgen tcp -i ens3f1 fe80::ee0d:9aff:fedb:11cd -p 1234
Connected to fe80::ee0d:9aff:fedb:11cd port 1234 from fe80::ee0d:9aff:fed8:f5d3 port 39500
[...]
XDP_DROP          3249504 pkts (    23878 pps)      4760015 KiB (   287 Mbits/s)
XDP_PASS           0 pkts (         0 pps)           0 KiB (         0 Mbits/s)
XDP_TX            516392331 pkts (  6217106 pps)    756434078 KiB ( 74605 Mbits/s) <--- retransmissions!
XDP_REDIRECT      545096864 pkts (  6534151 pps)    798481734 KiB ( 78410 Mbits/s)
```

Kernel samples tools

WiP: **moving** the tools from `kernel samples/bpf` into `xdp-tools`:

<https://github.com/xdp-project/xdp-tools/pull/158>

End: Questions?